

Introduction to Machine Learning

Definition of ML, Types of ML, Why ML?



Report an issue with this product

ISBN-10 ISBN-13 Edition

1782116613 978-1782116615 # Main

Other sellers on Amazon

New (16) from ₹400⁰⁰ **FREE** Delivery on orders over ₹499.

Frequently bought together

This item: The Brain ₹400⁰⁰

Behave : The bestselling exploration of why humans behave as th... ₹449⁹⁰

Livewired : The Inside Story of the Ever-Changing Brain ₹339⁹⁰


Total price: ₹1,188.00

Add all 3 to Cart

Some of these items are dispatched sooner than the others. Show details

Deals for you in Toys and Games See all deals

Page 1 of 6



youtube.com/watch?v=AUDHb-trnB0&t=1254s

kolmogorov arnold network

Symbolic formulas

Footnote 1: These datasets are quite contrived, just to benchmark the KAN scaling behaviour.

Footnote 2: We thank Yiping Lu for pointing out the our MLP baselines are not the best. He pointed us to a paper where two extra tricks are leveraged to achieve high precision in 1D & 2D cases, switching between Adam, LFGD and boosting. It is worthwhile to take this advanced setup and apply it to both MLPs and KANs, and see how these models compare.

KAN: Kolmogorov-Arnold Networks | Ziming Liu

Valence Labs 7.46K subscribers

28K views · 1 month ago · Granbe and Geometry Reading Group

Ready to play?

Decathlon's Running... Sponsored · decathlon.in

Power through deadlines

Supercharge your efficiency by handling multiple apps and assignments.

Sponsored · Apple

D-Flow: Differentiating through Flows for Controlled Generatio...

Valence Labs 1.1K views · 1 month ago

Definition of ML

- [Machine Learning is the] field of study that gives computers the ability to learn without being explicitly programmed.—Arthur Samuel, 1959
- A computer program is said to learn from experience E with respect to some task T and some performance measure P , if its performance on T , as measured by P , improves with experience E .—Tom Mitchell, 1997

Why to use ML?



- Machine Learning algorithm can often simplify code and perform better.
- Complex problems for which there is no good solution at all using a traditional approach: the best Machine Learning techniques can find a solution.
- Fluctuating environments: a Machine Learning system can adapt to new data.
- Getting insights about complex problems and large amounts of data.

Types of Machine Learning:

- Supervised learning
- Unsupervised learning
- Semisupervised learning
- Reinforcement Learning.



Supervised Learning

- The training data you feed to the algorithm includes the desired solutions, called labels.
- A typical supervised learning task is classification.
- Another typical task is to predict a target numeric value, such as the price of a car, given a set of features (mileage, age, brand, etc.) called predictors.



Classification Problem

- For the supervised problem the training data $D = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\}$ consists of inputs $x_i \in \mathbb{R}^d$ and output y_i pair.
- In the case of classification, the output set is categorical i.e. $y_i \in \{1, 2, \dots, K\}$
- If $y_i \in \{1, 0\}$, then the classification problem is called binary classification problem.
- If output set contains more than two class then it is called multi-class classification.



(a)



(b)




(c)

Regression Problem

- For regression problem the output or target variable is continuous.
- Input $x_i \in \mathbb{R}^d$ and the output $y_i \in \mathbb{R}$.
- Stock price prediction
- Temperature prediction


Unsupervised Learning

- The training data is unlabeled. The system tries to learn without a teacher.




Semi-supervised Learning

- Some algorithms can deal with partially labeled training data, usually a lot of unlabeled data and a little bit of labeled data. This is called semisupervised learning.



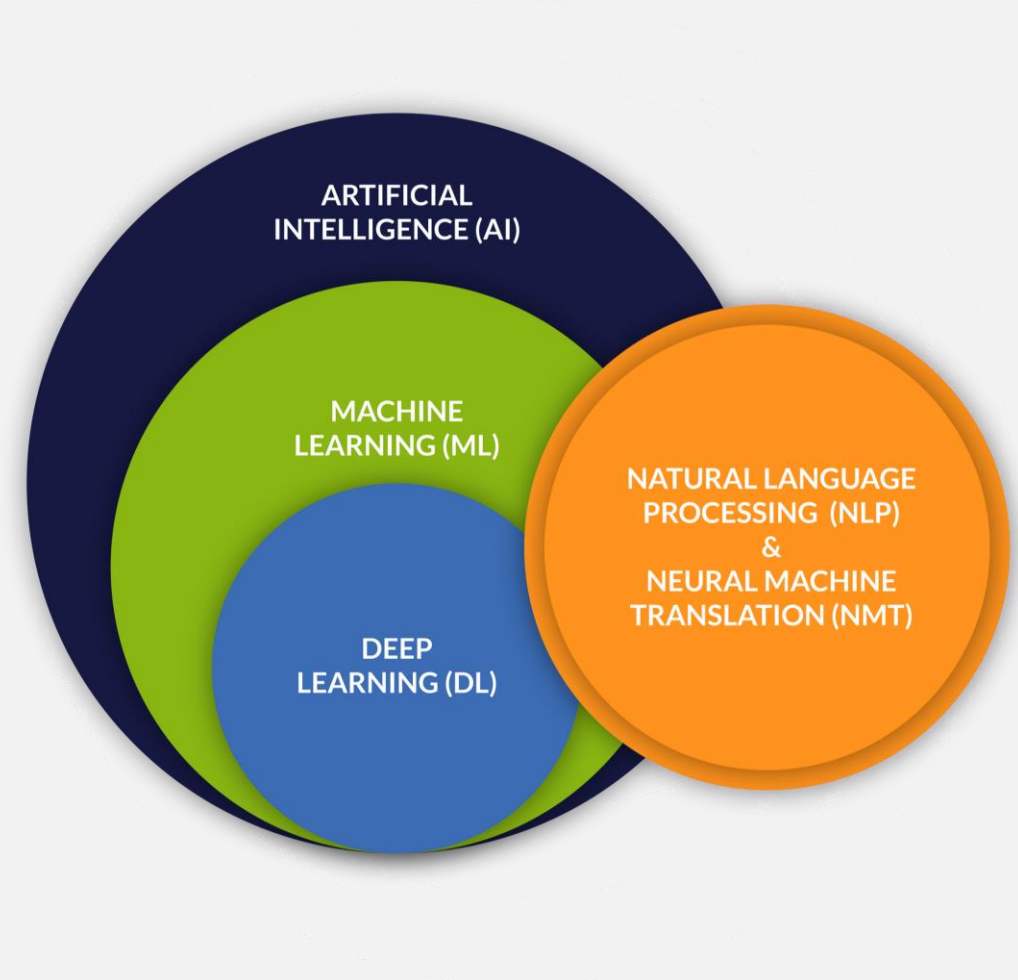
Reinforcement Learning

- The learning system, called an agent in this context, can observe the environment, select and perform actions, and get rewards in return.
- It must then learn by itself what is the best strategy, called a policy, to get the most reward over time.
- A policy defines what action the agent should choose when it is in a given situation.





AI, ML, and DL

- Artificial intelligence, or AI, is technology that enables computers and machines to simulate human intelligence and problem-solving capabilities.
- Machine Learning is a subset of artificial intelligence that helps you build AI-driven applications.
- Deep Learning is a subset of machine learning that uses vast volumes of data and complex algorithms to train a model.



Model as a function

- A predictor is a function that, when given a particular input example, produces an output. For simplicity we are considering here that the output is a single number i.e. a real number or scalar value.
- $f(x): \mathbb{R}^{N \times d} \rightarrow \mathbb{R}$
- $f(x) = \beta_0 + \beta x$



Model as Probability Distribution

- Machine learning techniques adopt a probabilistic approach, treating the target variable as a distribution of possible functions rather than a single function.
- $P(y|x) = P(y|f(x)) = P(y|x, \beta)$

Learning in Machine Learning

“A computer is said to learn from experience E with respect to some class of task T and performance measure P , if its performance at tasks in T , as measured by P , improves with experience E .”

Model Selection

- When we have a variety of models of different complexity, how should we pick the right one?
- The most conventional method is to select the model with minimum error.
- Train-Test-Validation Split
- Cross Validation



Reference

- Machine Learning-A Probabilistic Perspective-Kelvin P. Murphy
- Hands-on Machine Learning with Scikit-Learn, Keras and Tensorflow-Aurelien Geron